

Network stack personality in Android phone

Cristina Opriceana, **Hajime Tazaki** (IIJ Research Lab.)


Linux netdev 2.2, Seoul, Korea

07 Nov. 2017

Android

a platform of billions devices

- billions installed Linux kernel
- slow delivery of software updates
- Questions
 - When our upstreamed code available ?
 - What if I come up with a great protocol ?



<https://developer.android.com/about/dashboards/index.html>

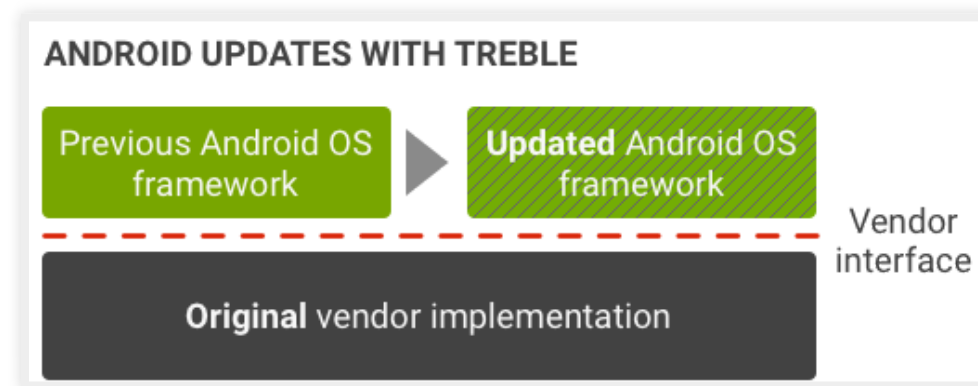
Android (cont'd)

- When our upstreamed code available ?
 - wait until base kernel is upgraded
 - backport specific function
- What if I come up with a great protocol ?
 - craft your own kernel and put into your image

Looong delivery to all billions devices

Approaches to alleviate the issue

- Virtualization (KVM on Android)
 - Overhead isn't negligible to embedded devices
- Project Treble (since Android O)
 - More modular platform implementation
- Fuchsia
 - Rewrite OS from scratch



<https://source.android.com/devices/architecture/treble>

An alternate approach

- network stack personality
 - use own network stack implemented in userspace
 - no need to replace host kernels
 - but (try to) preserve the application compatibility
- NUSE (network stack in userspace)
 - No delay of network stack update
 - Application can choose a network stack if needed

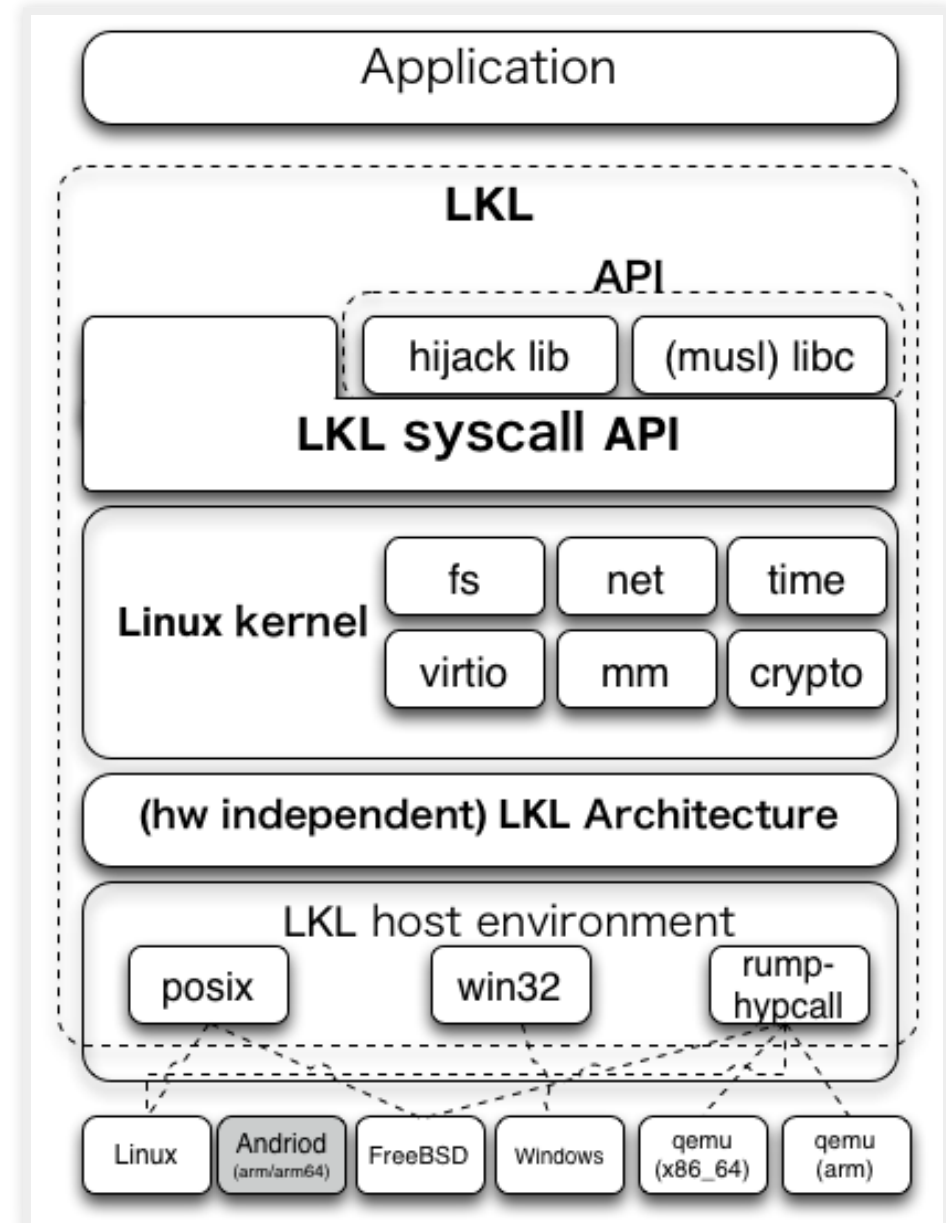
Userspace implementations

- Toys, Misguided People
- Selfish
- Don't care whether toy or not
- Do care whether it's useful or not

(TODO: replace image)

Linux Kernel Library intro (again)

- Out-of-tree architecture (h/w-independent)
- Run Linux code on various ways
 - with a reusable library
- h/w dependent layer
 - on Linux/Windows /FreeBSD uspace, unikernel, on UEFI,
 - (ns-3)
 - **Android**



LKL: current status

- Sent RFC (Nov. 2015)
 - no update on LKML since then
- have evolved a lot
 - various extensions (offload, SMP, json config, unikernel, etc)

<https://github.com/lkl/linux>

Extensions to LKL

- Android (arm/arm64) support (lkl/linux#372)
- raw socket extension (only handle ETH_P_IP) (not yet)
- hijack library enhance (not yet)
- HOWTO

```
% LD_PRELOAD=liblkl-hijack.so netperf XXX # console app  
% setprop wrap.app LD_PRELOAD=liblkl-hijack.so # Java app
```

hijack library

- For smooth replacement for Android UI app (java-based)
 - bionic is more familiar than glibc ?
 - only socket-related calls are redirected
 - mixture of host and lkl descriptors (can avoid)

TODO (a fig of hijack)

New feature introduction


- Example
 - Multipath TCP (<http://multipath-tcp.org/>)
 - out-of-tree for long time
 - verify site (cat /proc/net/mptcp base detection)
<http://amiusingmptcp.de/>

Demo

No penalty with userspace network stack ?


- Condition
 - To use Linux mptcp w/o replacing kernel
 - With tolerable amount of overhead
- Questions
 - Is NUSE working fine (Will users wanna use it) ?
 - How different from native Linux kernel ?

netperf measurement




- Client
 - Nexus5 android 6.01 (rooted)
 - LTE, wifi
 - LKL arm/android patched
 - or native kernel
- Server
 - Ubuntu 16.04 (amd64) on KVM
 - virtio/Etherlink (uplink: 100 Mbps)
 - mptcp-4.4.70 (v0.92)
- Software
 - netperf 2.7.x
 - 10 seconds TCP_STREAM, TCP_MAERTS
 - 5 trials, over 64-64K byte packet

Single path (Wi-Fi only, 9/14)



Tx (TCP_STREAM)

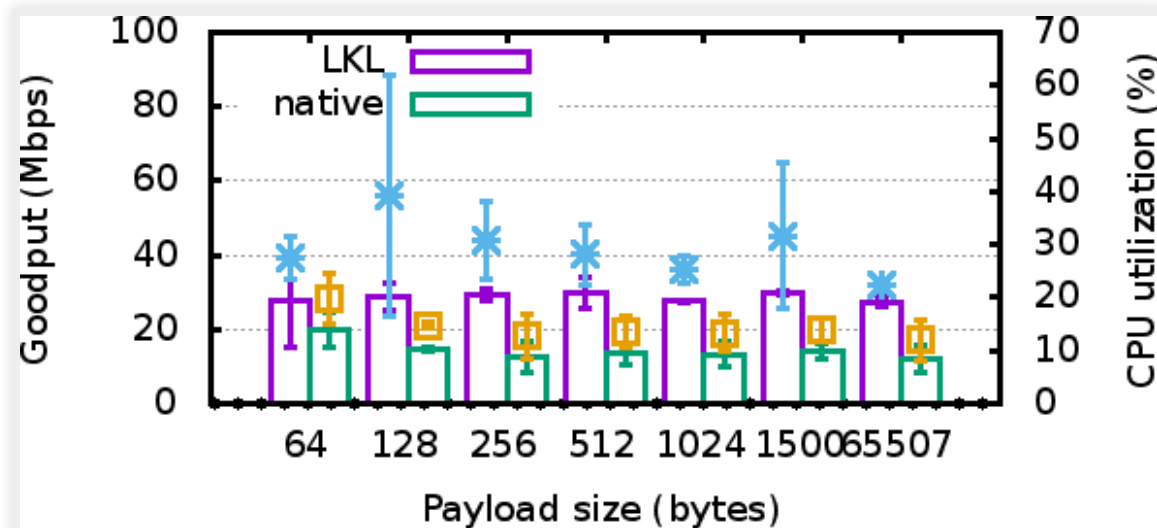
- Condition
 - phone: LKL v.s. native kernel



Rx (TCP_MAERTS)

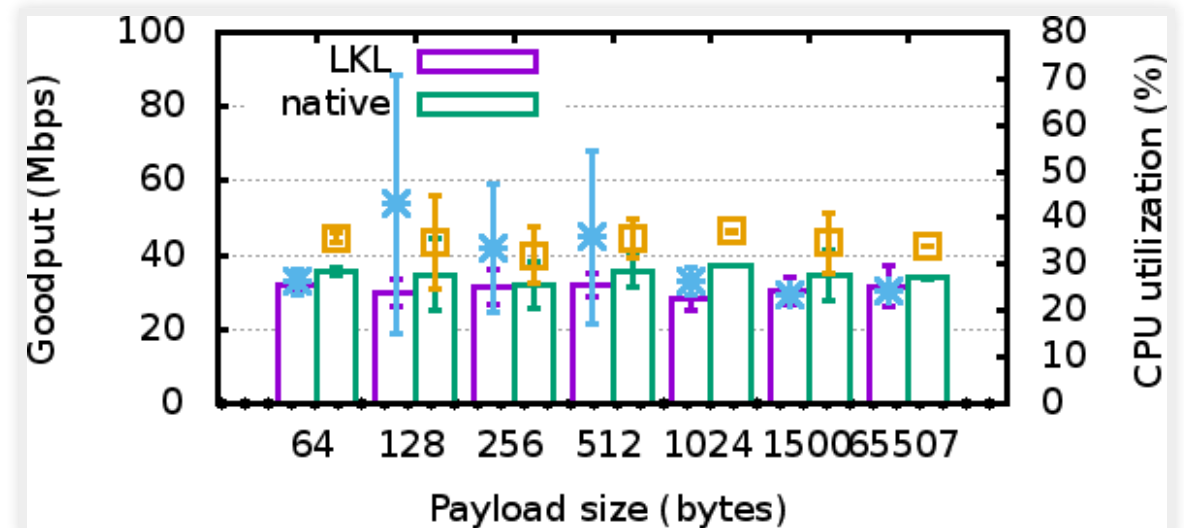
- Comparable goodput
- CPU utilization: LKL < native

Multipath TCP



Tx (TCP_STREAM)

- Condition
 - phone: LKL v.s. mptcp kernel



Rx (TCP_MAERTS)

- Tx with native kernel offers less goodput
 - even it's using multipath
- results are unstable

Multipath TCP (Korea/KT)

(TBA)

- Condition
 - phone: LKL v.s. mptcp kernel
- (TODO) To be measured

Observation

- IP conflicts may heavier
 - processed twice (host/lkl) per packet
- Results are always unstable
 - difficult measurement under wireless media

TODO (a fig of overhead)

Limitations

- Implementations
 - DHCP only boot time (handover will fail)
 - IPv4 only on cellular interface (rmnet0)
- Required tweaks
 - grant NET_RAW permission (packet socket)
 - need filter out RST packet from host

```
iptables -A OUTPUT -p tcp --tcp-flags RST RST -j DROP
```

Further investigations

- profiling
- other platform (iOS11 shipped userspace implementation)

Summary

- Use out-of-tree kernel as a library on Android
 - make your code easier to distribute
 - But with privileged installation/operation
- Comparable performance wireless communication